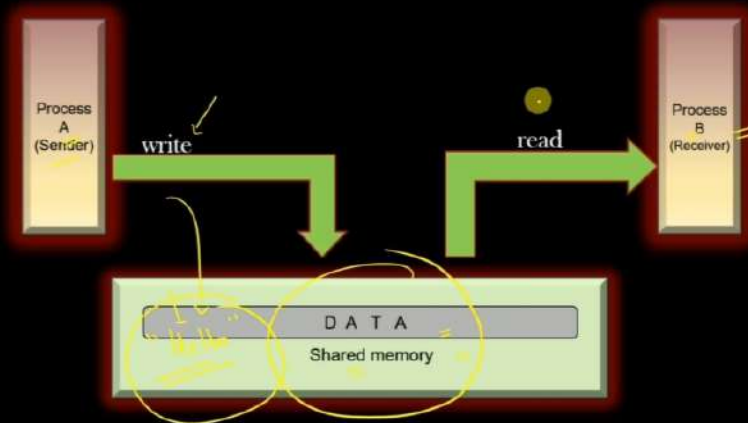


Shared Memory

1.1

IPC



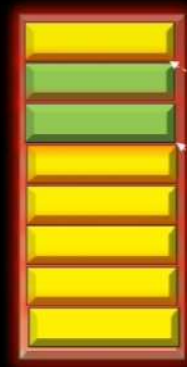
Memory Mapping

1.2

1. What is Memory Mapping ?
2. Goals ?
3. APIs to implement memory mappings
4. Shared Memory using Memory Mapping

Pre-requisite : Paging and Virtual Memory Basics

Codes : www.github.com/sachinities/IPC/NewSHM



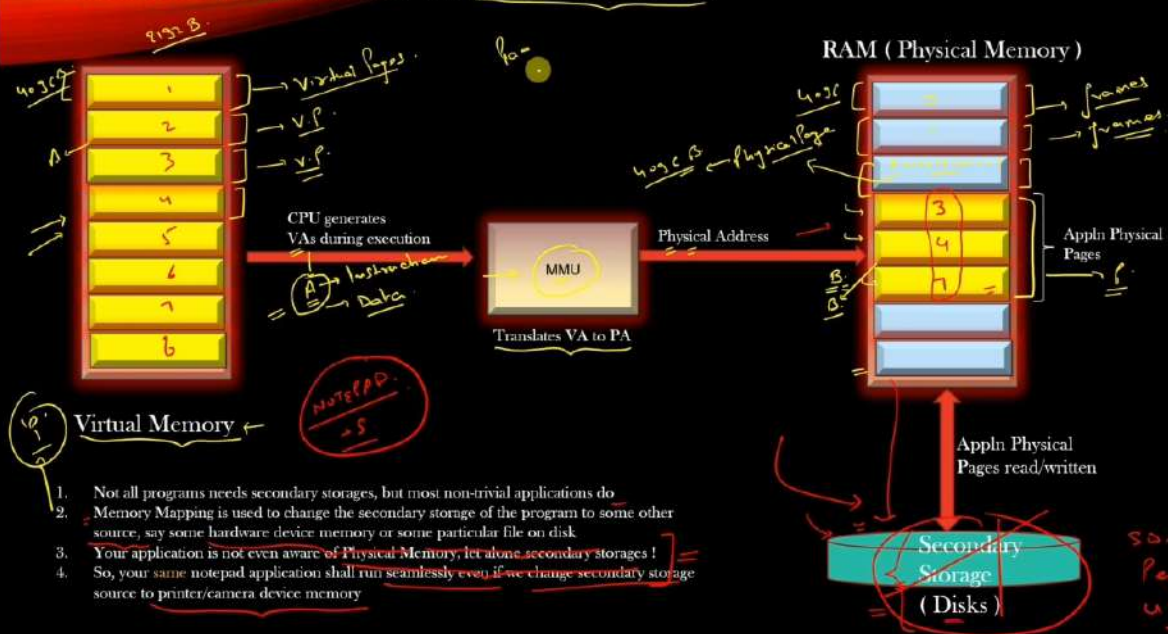


2.1

Paging
Recap
&
Memory Mapping

Memory Mappings

➤ Virtual Memory, Physical Memory and Secondary Memory setup



2.2

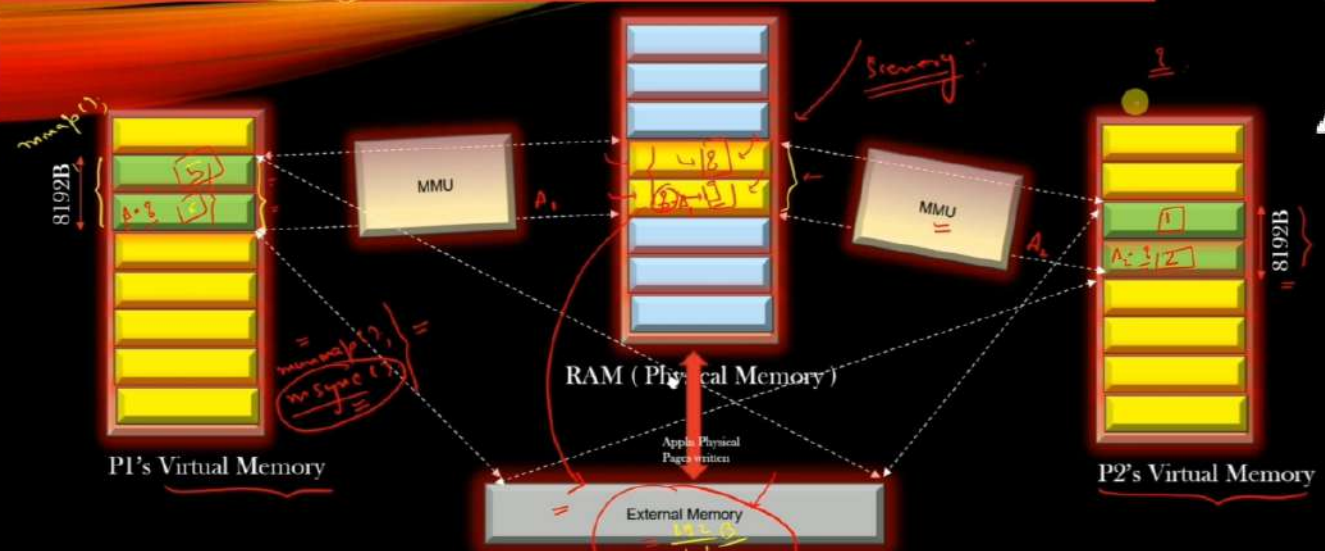
1. Not all programs need secondary storage, but most non-trivial applications do.
2. Memory Mapping is used to change the secondary storage of the program to some other source, say some hardware device memory or some particular file on disk.
3. Your application is not even aware of **Physical Memory**, let alone **secondary storage**!
4. So, your same notepad application shall run seamlessly even if we change secondary storage source to printer/camera device memory



4.1

Shared Memory

Memory Mappings → Shared Memory → Using External Data Source as Shared Memory



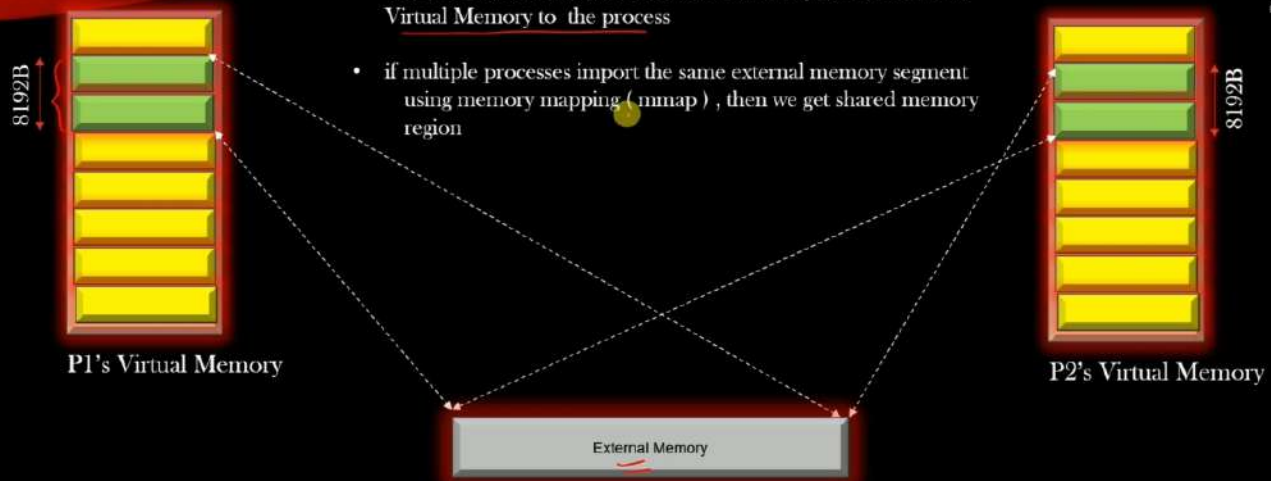
4.3

- Virtual Pages of both the Processes maps to same physical pages loaded in RAM
- Physical Pages in turn are read/written to external memory
- Rule : A process never can access any address outside of its VAS is never violated
- Any modification made by P1 in its shared VM, shall be seen by P2
- Let see logical view of this diagram on next slide

Memory Mappings → Shared Memory → Using External Data Source as Shared Memory

➤ Memory mapped Files - Logical View

- The end result is that that the external memory can be shown as Virtual Memory to the process
- if multiple processes import the same external memory segment using memory mapping (`mmap`), then we get shared memory region



4.4

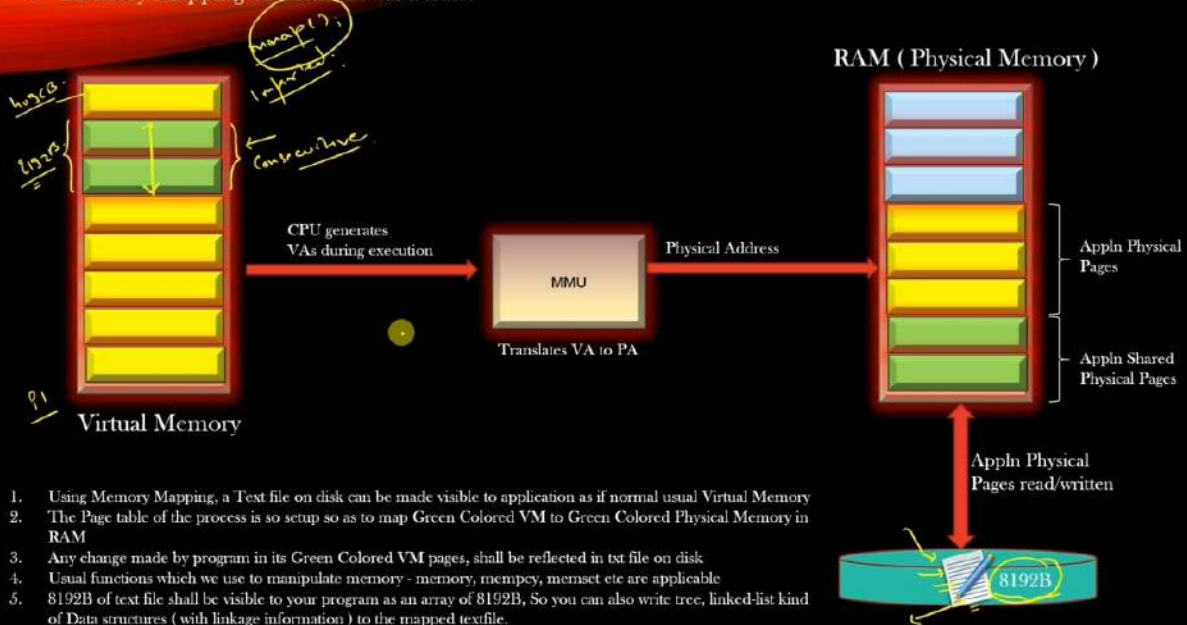


3.1

Memory Mapping Example

How Memory Mappings is Done in Linux - mmap ()

➤ Memory Mapping of a text file on a Disk



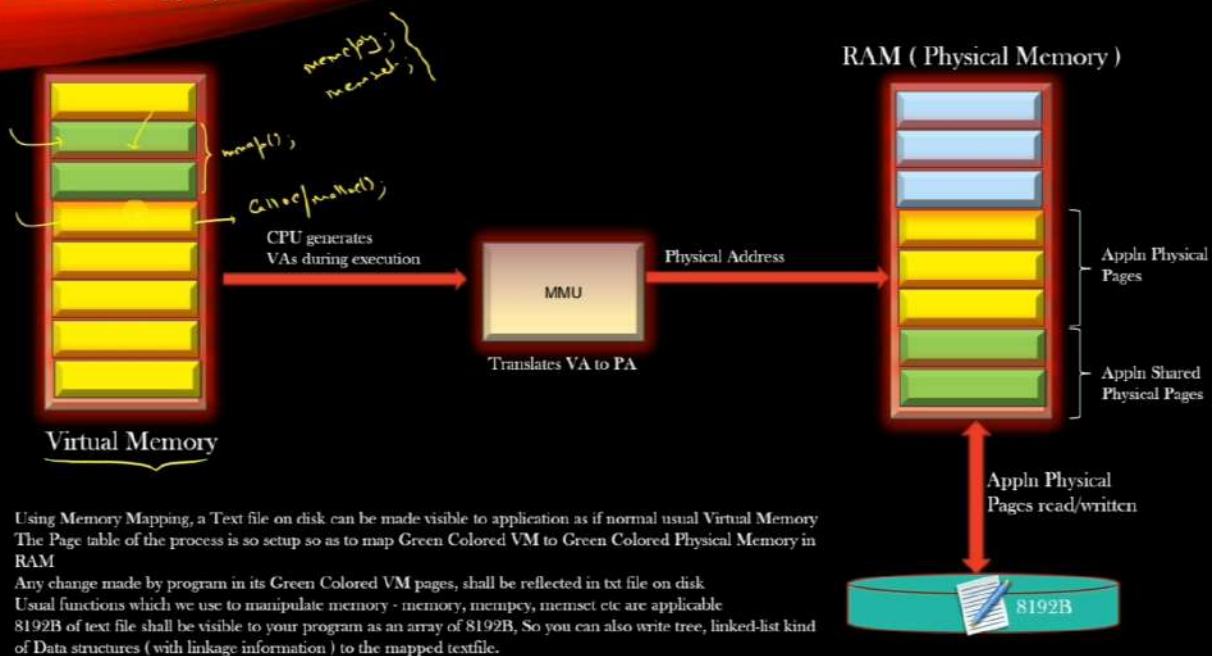
3.2

1. Using Memory Mapping, a Text file on disk can be made visible to application as if normal usual Virtual Memory
2. The Page table of the process is so setup so as to map Green Colored VM to Green Colored Physical Memory in RAM
3. Any change made by program in its Green Colored VM pages, shall be reflected in txt file on disk
4. Usual functions which we use to manipulate memory - memory, memcpy, memset etc are applicable
5. 8192B of text file shall be visible to your program as an array of 8192B, So you can also write tree, linked-list kind of Data structures (with linkage information) to the mapped textfile.

How Memory Mappings is Done in Linux - mmap ()

➤ Memory Mapping of a text file on a Disk

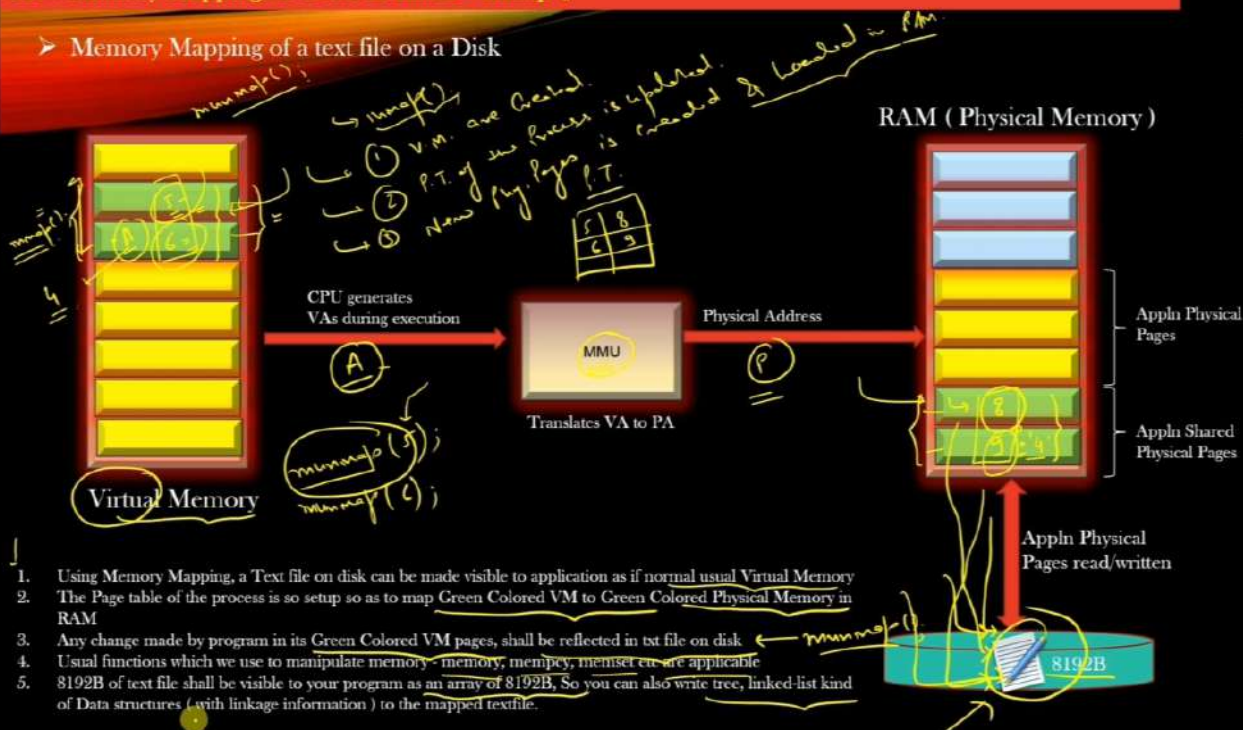
3.3



1. Using Memory Mapping, a Text file on disk can be made visible to application as if normal usual Virtual Memory
2. The Page table of the process is so setup so as to map Green Colored VM to Green Colored Physical Memory in RAM
3. Any change made by program in its Green Colored VM pages, shall be reflected in txt file on disk
4. Usual functions which we use to manipulate memory - `memory`, `memcpy`, `memset` etc are applicable
5. 8192B of text file shall be visible to your program as an array of 8192B, So you can also write tree, linked-list kind of Data structures (with linkage information) to the mapped textfile.

How Memory Mappings is Done in Linux - mmap ()

➤ Memory Mapping of a text file on a Disk

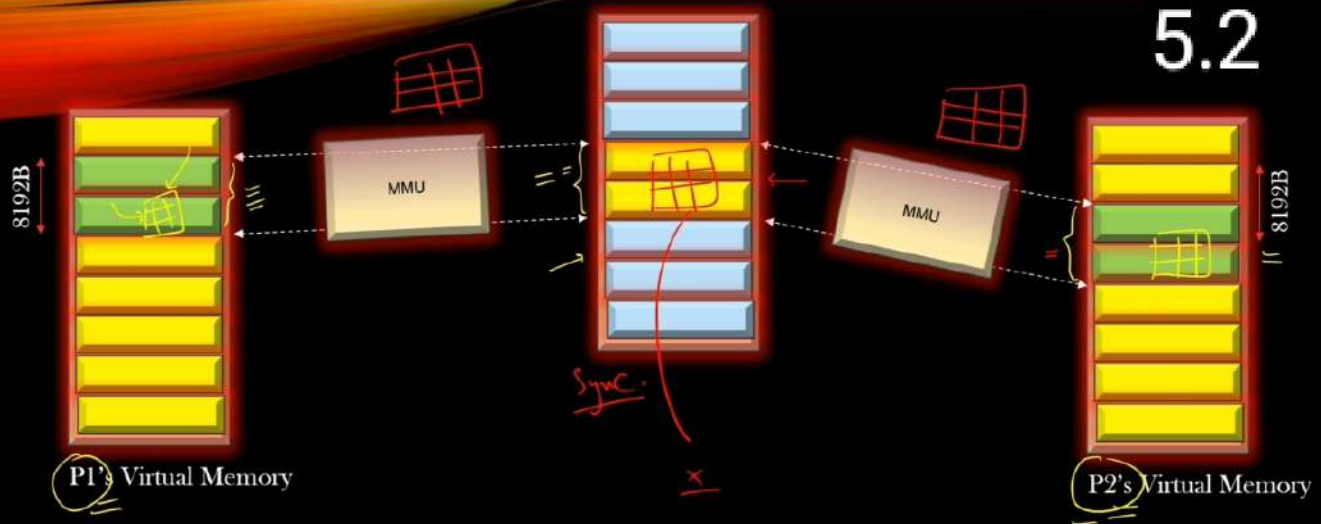


3.4



5.1

Using RAM
as
Shared Memory



- Virtual Pages of both the Processes maps to same physical pages loaded in RAM
- Used Widely for IPC
- Physical Pages in turn are read/written to external memory



6.1

Example Codes

➤ Example :

- Using `mmap()` as substitute for `malloc()`
- Using Text file as a data source
- Using RAM as a Data Source
- Memory Mapped being shared by Multiple Processes

Steps :

1. Initialize the Shared Memory segment
`shm_open()`
2. Define the size of the SHM segment -
`ftruncate()`
3. Map the Shared Memory segment to Data Source - `mmap()`
4. Use the shared Memory (read/write)
5. Destroy the mapping between process and Shared Memory segment - `munmap()`
6. Destroy shared memory segment -
`shm_unlink()`



7.1

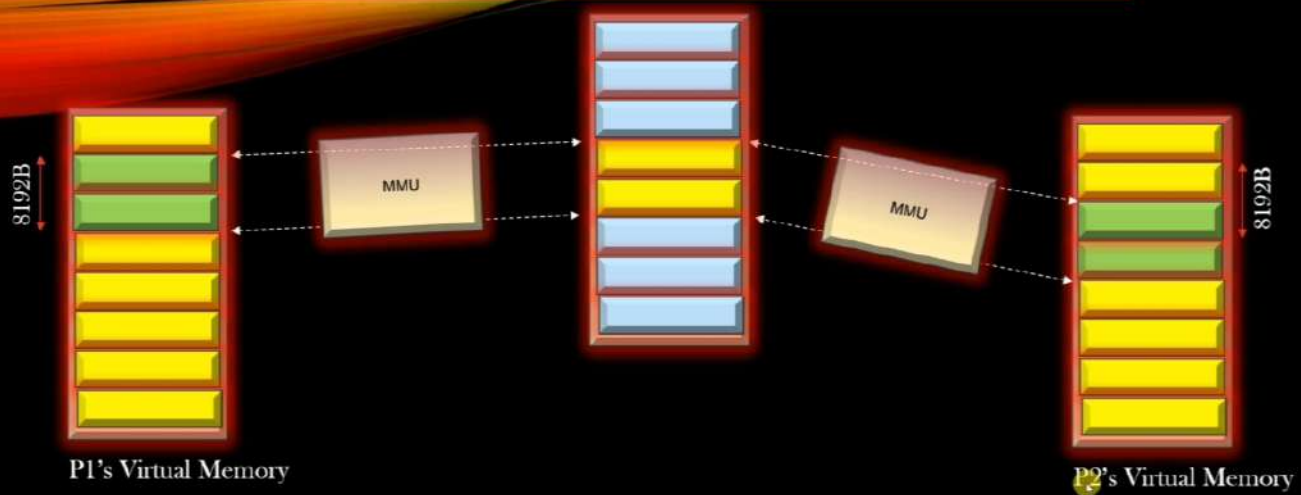
mmap()

```

C:\Program Files\Git\bin> cd C:\Program Files\Git\bin
C:\Program Files\Git\bin> gcc -o mmap_as_malloc_Demo1.exe mmap_as_malloc_Demo1.c
C:\Program Files\Git\bin> .\mmap_as_malloc_Demo1.exe
6 is no external data source or memory mapping involved. This example simply
7 shows how mmap() can be used as substitution for malloc()/free().
8 */
9
10 #if 0
11 void *
12 mmap(void *addr, size_t length, int prot, int flags,
13       int fd, off_t offset);
14
15 int munmap(void *addr, size_t length);
16 #endif
17
18 int
19 main(int argc, char **argv) {
20
21     int N=5;
22     int *ptr = mmap ( NULL, /* Let the OS choose the starting virtual address of the memory, just
23                        like you do not have a control as to what address malloc() will return */
24                      N*sizeof(int), /* Size of memory in bytes being requested for allocation */
25                      PROT_READ | /* Memory is Readable */
26                      PROT_WRITE, /* Memory is Writable */
27                      MAP_PRIVATE | /* This memory is not sharable with any other process, use MAP_SHARED instead */
28                      MAP_ANONYMOUS, /* This memory is not mapped to external data source, but only RAM by default */
29                      0, /* FD 0, since no external data source is specified */
30                      0 ); /* offset value as zero, since no external data source is specified */
31
32     if(ptr == MAP_FAILED){
33         printf("Mapping Failed\n");
34         return 1;
35     }
36
37     for(int i=0; i<N; i++)
38         ptr[i] = i*10;
39
40     for(int i=0; i<N; i++)
41         printf("%d\n", ptr[i]);
42 }

```

Memory Mappings → Shared Memory → Using RAM itself as Data Source



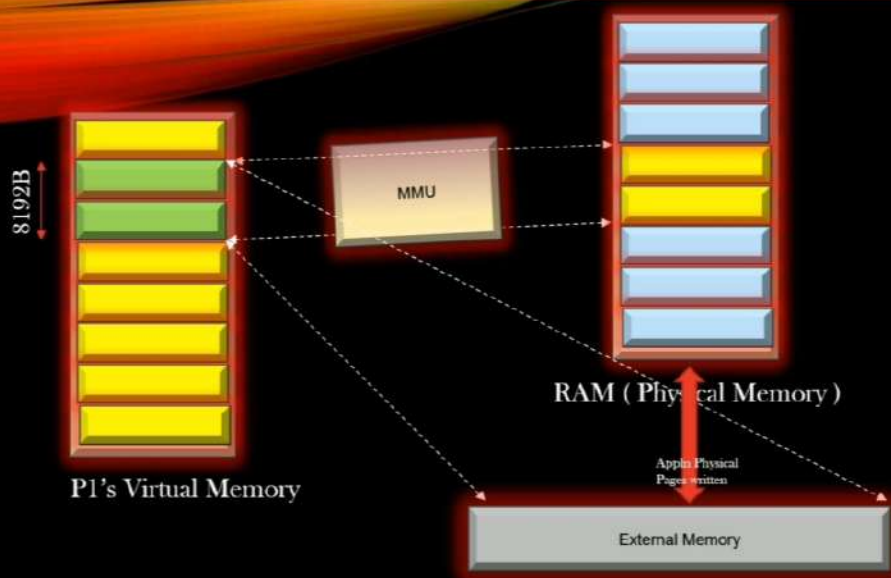
7.3

- Virtual Pages of both the Processes maps to same physical pages loaded in RAM
- Used Widely for IPC
- Physical Pages in turn are read/written to external memory



7.4

Mapping External Memory



Memory Mappings → Shared Memory → Using External Data Source as Shared Memory

```
0 #include <errno.h>
1 #include <unistd.h>
2 #include <memory.h>
3
4 /* This program shows how mmap() can be used to memory map the text
5  file present on disk into process's Virtual address space */
6 static void breakpoint () {}
7
8 typedef struct student_ {
9
10     int roll_no;
11     int marks;
12     char name[128];
13     char city[128];
14 } student_t;
15
16 int
17 main(int argc, char *argv[]) {
18     if(argc < 2) {
19         printf("File path not mentioned\n");
20         exit(0);
21     }
22     const char *filepath = argv[1];
23
24     /* Open the file in Read-Write OMode */
25     int fd = open(filepath, O_RDWR );
26
27     if(fd < 0) {
28         printf("%d\n", fd);
29         printf("could not open '%s'",
30                filepath);
31         exit(1);
32     }
33
34     /* Extract the size of the file */
```

7.6

Memory Mappings → Shared Memory → Using External Data Source as Shared Memory

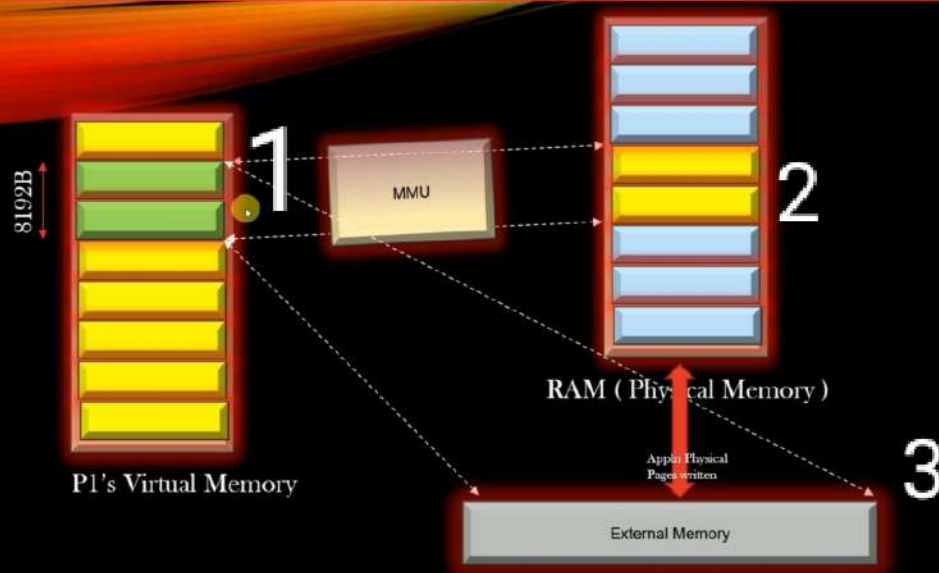
7.7

```
vm@ubuntu:~/src/IPC/NewSHM$ ls -l
total 140
-rw-rw-r-- 1 vm vm 539 Jul 5 14:02 file_demo2.txt
-rwxrwxrwx 1 vm vm 140 Jul 4 00:48 file_demo3.txt
-rw-rw-r-- 1 vm vm 1782 Jul 5 13:48 mmap_as_malloc_Demo1.c
-rwxrwxr-x 1 vm vm 19720 Jul 3 10:38 mmap_as_malloc_Demo1.exe
-rwxrwxr-x 1 vm vm 20936 Jul 4 00:51 mmap_file_read_Demo2.exe
-rw-rw-r-- 1 vm vm 2769 Jul 4 00:51 mmap_file_read_Demo3.c
-rw-rw-r-- 1 vm vm 2743 Jul 5 14:08 mmap_file_write_Demo2.c
-rwxrwxr-x 1 vm vm 21128 Jul 5 13:58 mmap_file_write_Demo2.exe
-rwxrwxr-x 1 vm vm 21096 Jul 4 00:48 mmap_file_write_Demo3.exe
-rw-rw-r-- 1 vm vm 964 Jul 4 00:19 test.c
-rwxrwxr-x 1 vm vm 18648 Jul 4 00:19 test.exe
-rw----- 1 vm vm 6 Jul 4 00:19 test_file
vm@ubuntu:~/src/IPC/NewSHM$ cat file_demo2.txt

vm@ubuntu:~/src/IPC/NewSHM$ ./mmap_file_write_Demo2.exe file_demo2.txt
vm@ubuntu:~/src/IPC/NewSHM$
```

Memory Mappings → Shared Memory → Using External Data Source as Shared Memory

The image shows a screenshot of the Visual Studio Code interface. On the left, the Explorer sidebar displays a project structure with folders like 'src' and 'test'. The main editor area shows a file named 'file_demo2.txt' with the text '1L, 539C' and '1L, 539C'. The bottom status bar indicates the file is at line 1, column 1, and the terminal shows the command 'file_demo2.txt' and the output '1L, 539C'.



7.9

- 1 . mmap() used for import external text file in virtual memory
- 2 . unmap() used to destroy virtual memory like ram space
- 3 . msync() used to copy the data from physical memory (ram) to text file

8.1

DESIGN CONSTRAINT

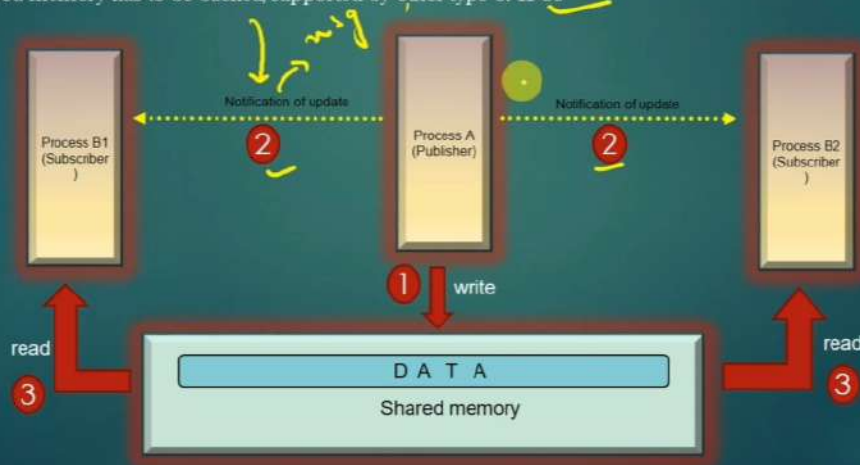
- Recommendation when it is good to use Shared Memory IPC :
 - Shared Memory approach of IPC is used in a scenario where :
 - Exactly one process is responsible to update the shared memory (publisher process)
 - Rest of the processes only read the shared memory (Subscriber processes)
 - The frequency of updating the shared memory by publisher process should not be very high
 - For example, publisher process update the shared memory when user configure something on the software
 - If multiple processes attempt to update the shared memory at the same time, then it leads to write-write conflict :
 - We need to handle this situation using Mutual Exclusion based Synchronization
 - Synchronization comes at the cost of performance
 - Because we put the threads to sleep (in addition to their natural CPU preemption) in order to prevent concurrent access to critical section

Design Constraints for using Shared Memory as IPC

8.3

- When publisher process update the shared memory :

- The subscribers would not know about this update
- Therefore, After updating the shared memory, publisher needs to send a notification to all ~~publishers~~ ^{Subscribers} which states that "shared memory has been updated"
- After receiving this notification, Subscribers can read the updated shared memory and update their internal data structures, if any
- The notification is just a small message, and can be sent out using other IPC mechanisms, such as Unix domain sockets Or Msg Queues
- Thus IPC using shared memory has to be backed/supported by other type of IPCs



9.1

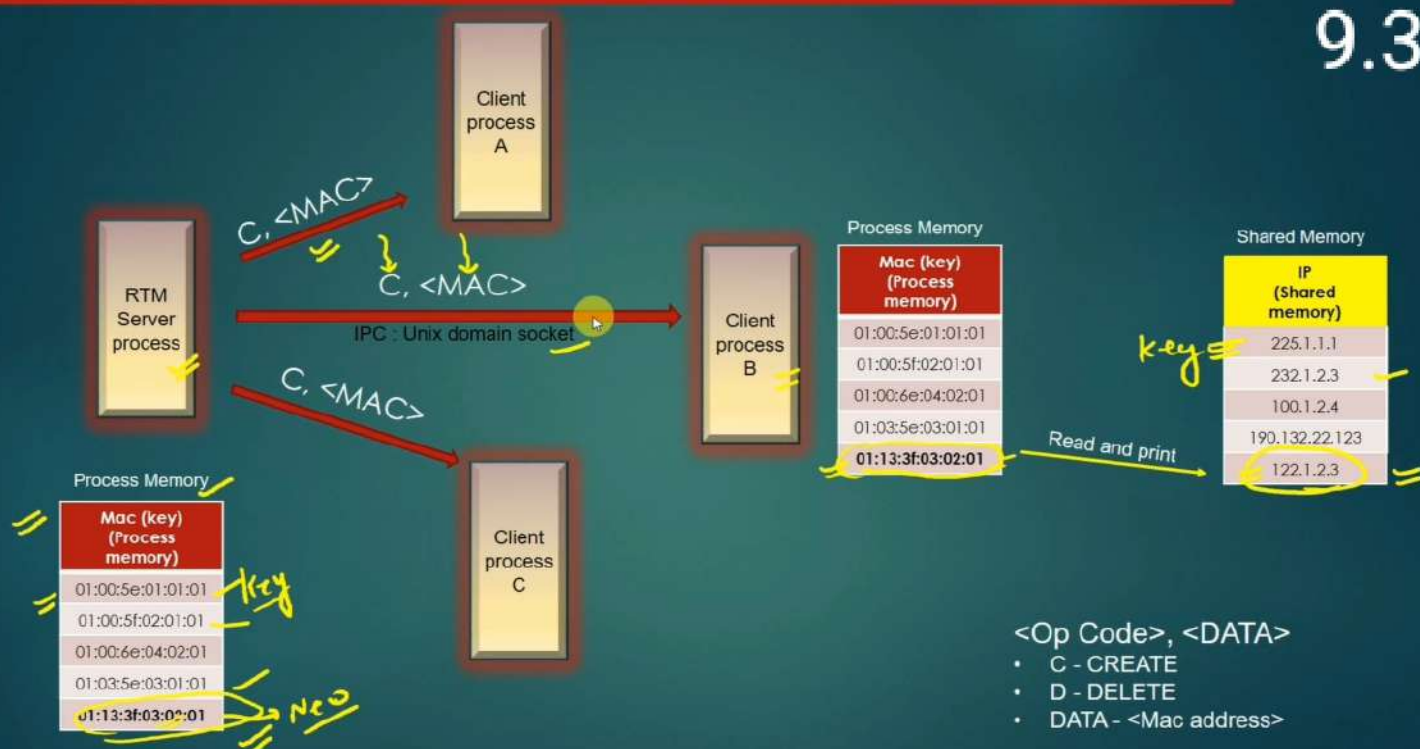
DATA SYNCHRONIZATION

SHARED MEMORY PROJECT

- We shall extend the project that you have done on Unix domain sockets
- In Unix domain socket project, the routing table manager server process synchronized routing table to all connected clients using Unix Domain sockets IPC
- In this project, the routing table manager server process also maintains another table called ARP table and it synchronizes it to all subscribed clients using Shared Memory IPC
- Let us discuss the project step by step. The functionality of Previous project will stay as it is, no changes

Project on shared Memory IPC

9.3



<Op Code>, <DATA>

- C - CREATE
- D - DELETE
- DATA - <Mac address>

Provide a menu-driven approach to show Mac table Contents on Server and Client processes

Project on shared Memory IPC

- All the processes – Server and clients stored only the shared memory key (the Mac address) in its internal data structure
- Server process adds the Data – the ip address in the shared memory corresponding to the mac address (key)
- Server process then syncs only the mac address (shm key) to rest of the connected clients using Unix domain sockets
- Client having received this new key (the mac address), stores the mac address in their private list. Using this mac address as key they can also access the corresponding ip address which was added by Server process in shared memory
- When the fresh clients get connected to server process, besides routing table, server also syncs the entire local mac address list to new client using unix domain socket

